

# AdaBits: Neural Network Quantization with Adaptive Bit-Widths

Qing Jin\*

Linjie Yang\*  
US AI Lab  
ByteDance Inc.

Zhenyu Liao

## Abstract

*Deep neural networks with adaptive configurations have gained increasing attention due to the instant and flexible deployment of these models on platforms with different resource budgets. In this paper, we investigate a novel option to achieve this goal by enabling adaptive bit-widths of weights and activations in the model. We first examine the benefits and challenges of training quantized model with adaptive bit-widths, and then experiment with several approaches including direct adaptation, progressive training and joint training. We discover that joint training is able to produce comparable performance on the adaptive model as individual models. We further propose a new technique named Switchable Clipping Level (S-CL) to further improve quantized models at the lowest bit-width. With our proposed techniques applied on a bunch of models including MobileNet-V1/V2 and ResNet-50, we demonstrate that bit-width of weights and activations is a new option for adaptively executable deep neural networks, offering a distinct opportunity for improved accuracy-efficiency trade-off as well as instant adaptation according to the platform constraints in real-world applications.*

## 1. Introduction

Recent development of deep learning enables application of deep neural networks across a wide range of platforms that present different resource constraints. For example, popular mobile apps such as TikTok and Snapchat on portable devices pose stringent requirements on response latency and energy consumption, while visual recognition system embedded in a self-driving vehicle is more demanding on fast and accurate prediction. The problem is more serious if other factors are taken into account, such as aging of hardware, battery conditions, as well as different versions of software systems. To serve applications under all these scenarios with drastically different requirements, different models tailored for different resource budgets can be devised either

manually [16, 17, 18, 31] or automatically through neural architecture search [43, 44, 33]. This strategy is beneficial for optimal trade-offs with a fixed combination of constraints, but is not economical, because it requires time-consuming training and benchmarking for each of these models, which prohibits instant adaptation to favor different scenarios. To tackle this problem, recent work focuses on training a single model that is flexible and scalable. For example, [38] proposes a method where the number of channels can be adjusted through changing the width-multiplier in each layer. Inspired by this work, [7] integrates adaptation of depth, width and kernel size altogether, and achieves better trade-offs between performance and efficiency through progressive training. [4] adopts the same strategy with scaling up factors, but uses simultaneous training algorithm to achieve improved predictive accuracy.

Surprisingly, albeit the above-mentioned methods achieve the desired flexibility of adaptive deployment, bit-width of weights and intermediate activations, as another degree of freedom, is almost overlooked in previous work. Suppose we can adaptively choose bit-width for a neural network during inference without further training, it will provide an distinct opportunity for more powerful model compression and acceleration. As an example, compared with model with full-precision, quantizing MobileNet-V2 to 6-bit compresses the model size by roughly  $4.74\times$  and reduces the BitOPs by  $14.25\times^1$ , while scaling the model's channel numbers by a width-multiplier of  $0.35\times$  only shrinks the model size by  $2.06\times$  and cuts down the FLOPs by  $5.10\times$ . Moreover, as presented in [3], 6-bit MobileNet-V2 demonstrates improved predictive capability than the full-precision counterpart, while reducing channel numbers to  $0.35\times$  will significantly impair its performance [31]. It becomes more contrastive if other constraints are taken into account, such as memory cost, latency and energy consumption. Additionally, adaptive bit-widths is generally applicable to most key building blocks of deep neural networks, including time-consuming convolutional and fully-connected layers. Mean-

\*Equal Contribution

<sup>1</sup>According to the IEEE Standard 754, floating-point number is represented with 23-bits mantissa, and here we simplify the analysis by approximating the effective BitOPs of floating-point multiplication with 23-bit fixed-point multiplication.

while, adaptive deployment will also introduce negligible computation, as discussed in [38]. Fig. 1 illustrates the basic concept of quantization with adaptive bit-widths.

At the first glance, adaptive bit-widths might be trivial and handy, as weights and activations with different precisions may not differ from each other very much. If so, model trained under some specific precision will be able to directly provide good performance under other bit-widths. However, as we will see in the following, such naïve method is not applicable, because important information will be lost during shrinkage or enlargement of bit-widths in the neural network. Even more deliberate method of progressively training quantized models with different bit-width(s) fails to achieve the optimal performance, as the finetuning process sabotages important property of the model, thus significantly diminishes the validation accuracy of the model when quantized back to the original bit-width.

All the above evidence indicates that quantization with adaptive bit-widths does not come as free lunch as it might appear, but is more subtle and involves new mechanisms that require meticulously designed techniques. In this work, we try to investigate this topic, and study specific methods to train quantized neural networks adaptive to different requirements. We utilize the state-of-the-art learning-based quantization method of Scale-Adjusted Training [3] as a baseline scheme for individual-precision quantization. We find that an adaptive model produced by a joint quantization approach with a key treatment to the clipping level parameters [9] is able to achieve comparable performance with individual-precision models on several bit-widths. The treatment to the clipping levels is named *Switchable Clipping Level (S-CL)*. S-CL accommodates large activation values for high-precision quantization, and prevents undesired increasing of clipping levels for low-precision cases. Through some empirical analysis, we find that unnecessarily large clipping levels might cause large quantization error, and impact the performance of quantized model, especially on the lowest precision. To our best knowledge, this work is the first to tackle this problem of producing quantized models with adaptive bit-widths.

This paper is organized as following. After summarizing some related works in Section 2, we first revisit the recent work of scale-adjusted training (SAT) [3], which is adopted in our whole study. In Section 4, we first illustrate potential benefits and challenges of quantization with adaptive bit-widths. Then we propose a joint training approach with a new technique named switchable clipping level based on the analysis of some baseline results. In Section 5, we show that with the proposed techniques, the adaptive models could achieve comparable accuracies as individual ones on different bit-widths and for a wide range of models including MobileNet-V1/V2 and ResNet-50.

## 2. Related Work

**Neural Network Quantization** Neural network quantization has long been studied since the very beginning of the recent blooming era of deep learning, including binarization [10, 11, 30, 5], quantization [39, 21, 41] and ensemble method [42]. Initially, uniform precision quantization is adopted inside the whole network, where all layers share the same bit-width [20, 40, 28, 25, 26, 37, 27, 19]. Recent work employs neural architecture search methods for model quantization, which implements mixed-precision strategy where different bit-widths are assigned to different layers or even channels [13, 36, 35, 34, 24]. [3] analyzes the problem of efficient training for neural network quantization, and proposes a scale-adjusted training (SAT) technique, achieving state-of-the-art performance. However, the possibility of developing a single model applicable at different bit-widths is still not well-examined, and it remains unclear how to achieve this purpose.

**Neural Architecture Search** Neural architecture search (NAS) gains increasing popularity since the early work of [43]. Later, the searching strategy is adopted in other aspects of optimizing neural networks, such as automatic tuning of various training hyper-parameters including activation function [29], data augmentation [12], and learning rate scheduler [1]. The NAS algorithms also benefit other tasks, such as generative adversarial networks [14], object detection [8] and segmentation [22]. As mentioned above, neural architecture search method for quantization is also actively studied in recent literature. However, NAS is computationally expensive, and usually requires time-consuming re-training or finetuning. Recent work has reduced the searching time by a large extent through one-shot architecture search [6, 32]. However, the resulting models are still inflexible, prohibiting their application in adaptive scenarios. Generally, conventional NAS methods are more suitable for optimizing a single model under specific resource constraints.

**Adaptive neural networks** Different from but related to NAS, [38] proposes to simultaneously train a single model with different width multipliers, to achieve instant adaptation for different application requirements. Following this line, [7] explores adjustment of width, depth and kernel sizes simultaneously, achieving better predictive accuracy under the same computational constraints through progressive training. [4] extends similar strategy to large-size models, and further employs a NAS algorithm to discover better models. However, these methods neglect the option of quantization with different bit-widths in their strategies, leaving quantization with adaptive bit-widths an open problem.

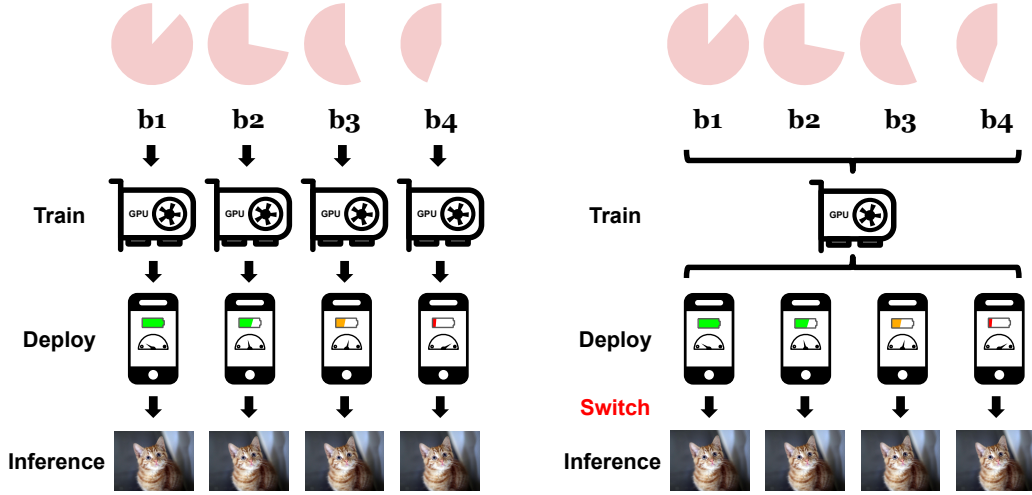


Figure 1. Deployment of neural networks with different bit-widths according to the computational budget. Left: Individually train several quantized models with different bit-widths for each scenario. Right: Train a single model quantized with adaptive bit-widths and switch to the proper bit-width in real application based on the device condition.

### 3. Revisiting Scale-Adjusted Training (SAT)

Quantization usually comes with performance degeneration, as the model capacity is significantly reduced in comparison with the full-precision counterpart. However, a recent study [3] demonstrates that a large portion of accuracy degradation is caused by inefficient training where learning-based quantization, potentially acting as a regularization, actually provides an opportunity to improve generalization capability. The key idea is that quantized models usually enforce large variance in their weights, which brings about over-fitting issue during training. Based on this finding, [3] proposes a simple yet effective method, called scale-adjusted training (SAT), which scales the weights down to a healthy level for network optimization. Specifically, constant scaling is applied to the quantized weights of linear layers without BN by

$$Q_{ij}^* = \frac{1}{\sqrt{n_{\text{out}} \text{VAR}[Q_{rs}]}} Q_{ij} \quad (1)$$

where  $Q_{ij}$  is a quantized weight and  $n_{\text{out}}$  is the number of output neurons in this layer. By combining with a quantization approach named parameterized clipping activation (PACT) [9], SAT facilitates more efficient training, enabling quantized models to perform consistently and significantly better than conventional quantization techniques, sometimes even surpassing their full-precision counterparts.

Due to the numerous algorithms for neural network quantization, it is difficult, if not impossible, to experiment with different quantization algorithms for the adaptive bit-widths problem. To this end, we adopt the PACT algorithm with the SAT technique, which gives the state-of-the-art perfor-

mance for neural network quantization, throughout all of our experiments. For brevity, we refer to this approach as SAT.

### 4. Quantization with Adaptive Bit-widths

In this section, we first examine the benefits and challenges of quantization with adaptive bit-widths. We explore direct adaptation and progressive quantization as two straightforward methods towards this goal with unsatisfying results. We then propose a novel joint quantization approach to deal with the challenge and achieve the same level performance with the adaptive models compared to the individual models.

#### 4.1. Benefit and Challenges

Neural network quantization provides significant reduction in model size, latency, and energy consumption. The SAT approach shows comparable accuracy of quantized model with their full-precision counterparts [3]. This seems to reduce the necessity to develop quantized models with adaptive bit-widths. However, we argue that quantized models with different bit-widths still pose a great opportunity for flexible and adaptive deployment since models with larger bit-width are still consistently better than those with smaller bit-width. We show the validation accuracy of quantized MobileNets produced by SAT [3, 18, 31] on ImageNet dataset with respect to the model sizes and BitOPs in Fig. 2. According to Fig. 2, there is a noticeable trade-off between accuracy and efficiency on quantized models. In the following, we will first investigate two straightforward methods for adaptive bit-widths, which will reveal some key challenges of this problem.




Figure 2. Trade-offs between prediction accuracy and model sizes or BitOPs on quantized MobileNet V1/V2 on ImageNet dataset.

#### 4.1.1 Modified DoReFa Scheme

Before more detailed analysis, we would like to emphasize a distinct difficulty encountered in quantized models with adaptive bit-widths. The DoReFa scheme [39] is adopted in the original SAT method for weight quantization, where weights are quantized with

$$q_k(x) = \frac{1}{a} \left[ ax \right] \quad (2)$$

Here,  $\left[ \cdot \right]$  indicates rounding to the nearest integer, and  $a$  equals  $2^b - 1$  where  $b$  is the number of quantization bits. However, such a scheme is not practical for quantization with adaptive bit-width, as there is no direct mapping between weights quantized to different bit-widths, disabling direct conversion of quantized models from a bit-width to lower bit-widths. It necessitates storage of the full-precision weights, and the quantization procedure needs to be repeated for different bit-widths during model deployment. This significantly increases the size of the stored model, and greatly limits the applications of the model. To accommodate simple conversion of quantized models, we modify the DoReFa scheme to use a quantization function given by

$$q_k(x) = \frac{1}{\hat{a}} \min \left( \left\lfloor \hat{a}x \right\rfloor, \hat{a} - 1 \right) \quad (3)$$

Here,  $\left\lfloor \cdot \right\rfloor$  indicates the floor rounding function, and  $\hat{a}$  equals  $2^b$  where  $b$  is the number of quantization bits. This quantizing function does not differ quite much from that for the original DoReFa scheme, and should give similar performance for quantized model. Moreover, it enables direct adaptation from higher bit-width to lower bit-width through discarding lower bits in the weights directly. In the following, we utilize this scheme to explore quantization with adaptive bit-widths, and directly borrow baseline results from the SAT paper [3] for comparison, assuming that such minor modification does not affect model accuracy.

#### 4.1.2 Direct Adaptation

We first investigate whether quantized models trained on one bit-width can be directly used on other bit-widths. This cheap approach could be viable since weights with different bit-widths might be close to each other in value. To check if this method is practical, we evaluate the validation accuracy of ResNet-50 on ImageNet by adjusting the bit-width to several different settings, where the original weights are trained under either the lowest or the highest bit-width (2 bit and 4 bit in this case, respectively). As indicated in previous research [3], quantization with different bit-widths entails difference in variances of weights and activations, as illustrated in Fig. 4.1.2. Thus the networks trained on one bit-width suffer from the mismatch of the layer statistics when evaluated on another bit-width. To alleviate this problem, we apply batch norm (BN) calibration introduced in [2] to calibrate the statistics in batch normalization layers for reasonable comparison.

The results with and without BN calibration are listed in Table 1, together with performance of models trained and evaluated under the same bit-width using SAT. It is shown that without BN calibration, models trained on one bit degenerate significantly on another bit. With BN calibration, model trained on 2 bit successfully preserves the performance on larger bits, but is still inferior to results achieved by directly training on the large bits; model trained on 4 bit still degenerates severely on smaller bits. In summary, models trained and evaluated in different bit-widths are not suitable for adaptive deployment of quantization models due to the difference in training and evaluation settings. Particularly, models trained with larger bit-width suffer from more serious performance degeneration when quantized to lower precision, while training with smaller bit-width limits the potential of models deployed on higher precision.

#### 4.1.3 Progressive Quantization

The above analysis demonstrates that quantization with adaptive bit-widths is not directly available from models




Figure 3. Impact of quantization on variances of weights and activation under different bit-widths. The variance of both quantized weights and activations gets larger when the bit-width gets smaller.

Model	4 bit	3 bit	2 bit
4 bit Trained (wo/ BN calib)	76.3	67.1	0.3
2 bit Trained (wo/ BN calib)	41.1	48.7	73.3
4 bit Trained (w/ BN calib)	76.3	73.2	20.3
2 bit Trained (w/ BN calib)	73.4	73.2	73.3
SAT [3]	76.3	75.9	73.3

Table 1. Direct adaptation of models trained on 2 and 4 bits on different bit-widths, with and without batch norm calibration. Results are top-1 validation accuracy (%) of ResNet-50 on ImageNet.

trained with individual bit-width. In this section, we investigate the possibility of progressive training, where a quantized model is trained on multiple bit-widths sequentially. It can be conducted in two ways, where the bit-width for training can be gradually increased or decreased. We experiment with both methods for ResNet-50 on ImageNet. For either case, we use the model trained individually with the highest (lowest) bit-width as the initial point, which is finetuned with the second highest (lowest) bit-width, and further finetuned with the next bit-width. We continue this finetuning process until all bit-widths under consideration are processed. For each phase of finetuning, the same hyper-parameters are adopted as those for training individual quantization. Also, BN calibration is applied to the final model on different bit-widths for reasonable comparison. The results are shown in Table 2.

In Table 2, the model first trained with 2 bit and finetuned with ascending bit-width achieves good result at the final 4 bit, but is corrupted on lower bits. The model first trained with 4 bit and fine-tuned with descending bit-width only achieve slightly better performance than directly applying 2 bit model on multiple bits in Table 1, which does not preserve performance of the higher 3 and 4 bits. The above results indicate that progressive training might have introduced undesired perturbation to the model trained previously, which impairs its original performance. This shows the progressive training method is still not suitable for models with adaptive bit-widths.

Model	4 bit	3 bit	2 bit
Ascending Bit-width	76.3	73.4	29.5
Descending Bit-width	73.9	73.6	73.5
SAT [3]	76.3	75.9	73.3

Table 2. Results of progressive quantization with ascending/descending bit-widths of ResNet-50 on ImageNet. Results are top-1 validation accuracy (%).

Model	8 bit	6 bit	5 bit	4 bit
Vanilla AdaBits	72.4	72.5	72.1	70.8
SAT [3]	72.6	72.3	71.9	71.3

Table 3. Results of Vanilla AdaBits with MobileNet-V1 on ImageNet with four bit-widths. Results are top-1 validation accuracy (%).

## 4.2. Joint Quantization

The above results shows that sequential training does not preserve the model characteristics in previously trained bit-widths, which indicates that the model weights for different bit-width should be jointly optimized. Specifically, we adopt a joint training approach similar to slimmable neural networks [38]. Instead of training models with different channel numbers, we simultaneously train models under different bit-widths with shared weights. Also, as mentioned above, quantization with different bit-widths leads to different variances of quantized weights and activations. Based on this, we adopt the switchable batch normalization technique introduced in [38]. We call this method *Vanilla AdaBits*, and the performance is listed in Table 3. It can be seen that models trained with all of the bit-widths achieve comparable performance as those trained individually, which validates the effectiveness of this approach. However, there is still a performance gap for the lowest bit-width at 4 bit, which is 0.5% lower than the individually trained model. This is undesired and further improvement needs to be made.

In the PACT algorithm adopted by SAT, the activations of each layer will be first clipped by a learned parameter  $\alpha$






Figure 4. Clipping levels in different layers for models trained individually with different bit-widths (solid lines) or trained with Vanilla AdaBits (dashed line). Note that the clipping level of a layer refers to the clipping level for the output of this layer. The outputs of the last layer are not clipped.

named *clipping level*, and then quantized to discrete numbers. Specifically, an activation value  $x$  is first clipped to the interval  $[0, \alpha]$ , then scaled, quantized and rescaled to produce the quantized value  $q$  as

$$\tilde{x} = \frac{1}{2} \left[ |x| - |x - \alpha| + \alpha \right] \quad (4a)$$

$$q = \alpha q_k \left( \frac{\tilde{x}}{\alpha} \right) \quad (4b)$$

Note that in the original paper of PACT [9], the authors found that different bit-widths result in different clipping levels. In the Vanilla AdaBits, the clipping levels of different bit-widths are shared, which may potentially disturb the optimization process of the network.

To understand the underlying mechanism of the degeneration at the lowest bit-width, we plot the clipping levels from different layers in models trained individually with different bit-widths. As shown in Fig. 4, the clipping levels strongly correlates with the bit-width. For the individually trained models, higher bit-widths result in larger values of clipping levels. In the Vanilla Adabits model, the learned clipping levels tend to be smaller than those of high-precision cases, but are larger than those from the model with the lowest bit-width. To understand the relationship between quantization error and clipping levels, we study the characteristics using a synthetic linear layer with 1000 input neurons where the weights are sampled from  $\mathcal{N}(0, 1/1000)$  and activations are sampled from a uniform distribution on the interval  $[0, 1]$ . For each bit-width, the products of the weights and the activations are fed to the ReLU function to obtain quantized outputs with different clipping levels. The relative error between the full-precision outputs and the quantized outputs are calculated. We plot this quantization error with respect to the clipping levels in Fig. 5. It shows that different bit-widths have different behaviors. Quantization error only increases slowly with increase of clipping level for higher

bit-width while it increases significantly with increase of clipping level for lower bit-width. Based on the results from Fig. 4, the clipping levels learned by Vanilla Adabits may substantially increase the quantization error for the lowest 4 bit, but do not affect those of the other bit-widths much. Note this is only a qualitative analysis and the quantization errors shown in Fig. 5 are not proportional to the quantization errors in the trained networks.




Figure 5. Relationship between relative quantization error and the clipping levels  $\alpha$  for different bit-widths with a synthetic layer. The dots denote the optimal values of  $\alpha$  for least quantization error at different bit-widths.

#### 4.2.1 Switchable Clipping Level

The above observation indicates that facilitating proper clipping levels for each bit-width could be a key factor for optimal performance of AdaBits models. One set of shared clipping levels is difficult, if not impossible, to satisfy requirements from different bit-widths. To this end, we propose a simple treatment to the clipping levels, named *Switchable Clipping Level (S-CL)*, that employs independent clipping levels for different bit-widths in each layer. During training of quantized models with adaptive bit-widths, S-CL switch to a corresponding set of clipping levels for each bit-width in all layers. This avoids the clipping level parameters being interfered by other bit-widths, especially undesired quantization error introduced by too large or too small clipping levels for the current bit-width. In this way, the performance degeneration issue on the lowest bit-width using Vanilla Adabits can be alleviated.

The model size is almost unchanged with S-CL, which is a negligible portion of less than 0.1%. For instance, the byte size ratio of clipping level with other trainable parameters is 0.0246% for MobileNet-V1, 0.0588% for MobileNet-V2, and 0.0084% for ResNet-50. Meanwhile, S-CL introduces almost no runtime overhead. After re-configuring the model with desired bit-width, it becomes a normal network to run without additional latency and memory cost. These advantages make it a very practical and economical solution to the adaptive bit-widths problem.

---

**Algorithm 1** AdaBits: Training quantized neural network  $M$  with adaptive bit-width.

---

**Require:** Define *switchable bit-width list* for  $M$  with adaptive bit-width, for example,  $[8, 6, 5, 4]$  bits.

- 1: Initialize shared convolutions and fully-connected layers for  $M$ .
  - 2: Initialize independent batch normalization parameters for each *bit-width* in *switchable bit-width list*.
  - 3: Initialize independent clipping levels parameters for each *bit-width* in *switchable bit-width list*.
  - 4: **for**  $i = 1, \dots, n_{iters}$  **do**
  - 5:   Get next mini-batch of data  $x$  and label  $y$ .
  - 6:   Clear gradients of weights,  $optimizer.zero\_grad()$
  - 7:   **for** *bit-width* in *switchable bit-width list* **do**
  - 8:     Switch the batch normalization parameters and clipping levels of current bit-width on network  $M$ .
  - 9:     Execute sub-network at current bit-width,  $\hat{y} = M'(x)$
  - 10:     Compute loss,  $loss = criterion(\hat{y}, y)$ .
  - 11:     Compute gradients,  $loss.backward()$ .
  - 12:   **end for**
  - 13:   Update weights,  $optimizer.step()$ .
  - 14: **end for**
- 

#### 4.2.2 Details of Training Procedure

Our final approach to train quantized neural network with adaptive bit-widths is straight-forward. The losses are directly summed over all bit-widths in forward pass and back-propagated for each bit-width. During training, weights are updated after accumulating gradients using all bit-widths, and we do not tune hyper-parameters specifically. Our final algorithm with the S-CL approach is named *AdaBits*. The algorithm is summarized in Algorithm 1, where *switchable bit-width list* is the predefined list of bit-widths for quantization.

## 5. Experiments

We evaluate our AdaBits algorithm on the ImageNet classification task and compare the resulted models with those quantized individually with different bit-widths. After that, we analyze the clipping levels in different layers from an AdaBits model. Finally, we give discussion and present some future work.

### 5.1. ImageNet Classification

To examine our proposed methods, we quantize several representative models with adaptive bit-widths, including MobileNet-V1/V2 and ResNet-50, and evaluate them on the ImageNet dataset, using AdaBits algorithm.

We follow the same quantization strategy as SAT [3], which first trains a full-precision model, and then uses it as

initialization for training the quantized model. The same training hyperparameters and settings are shared between pretraining and finetuning, including initial learning rate, learning rate scheduler, weight decay, the number of epochs, optimizer, batch size, etc. The input images to the model are set to unsigned 8bit integer (uint8), and no standardization (neither demeaning nor normalization) is applied. For the first and last layers, weights are quantized with bit-width of 8bit [9], while the input to the last layer is quantized with the same precision as other layers. Meanwhile, bias in the last fully-connected layer(s) and the batch normalization layers are not quantized.

To make a fair comparison, we adopt the same hyperparameters as SAT [3]. The learning rate is initialized to 0.05, and updated every iteration for totally 150 epochs with a cosine learning rate scheduler [23] without restart. Parameters are updated by a SGD optimizer, Nesterov momentum with a momentum weight of 0.9 without damping. Weight decay is set to  $4 \times 10^{-5}$ . For MobileNet-V1/V2, the batch size is set to 2048, while for ResNet-50 it is 1024. The warmup strategy suggested in [15] is adopted by linearly increasing the learning rate every iteration to a larger value (batch size/256  $\times$  0.05) for the first five epochs before using the cosine annealing scheduler. The input image is randomly cropped to  $224 \times 224$  and randomly flipped horizontally, and is kept as 8 bit unsigned integer with no standardization applied. Besides, we use full-precision models with clamped weight as initial points to finetune quantized models.

The results for these models are summarized in Table 4, where we list the Top-1 accuracy on ImageNet classification task, together with model size and BitOPs. We use a prefix of AB- to indicate model quantized with AdaBits. Result of the SAT approach [3] is also reported as reference, which based on our knowledge presents the state-of-the-art performance on model quantization. We find that our method is able to achieve almost the same performance as individual quantization for all models with all bit-widths. Compared to progressive quantization with ascending bit-width in Table 2, AdaBits approach on ResNet50 significantly boost the performance on lowest 2 bit. Compared to progressive quantization with descending bit-width, AdaBits boost accuracy of 2.2% on 4 bit and 2.2% on 3 bit on ResNet50. Compared to Vanilla AdaBits, our final approach with S-CL increase performance on the lowest 4 bit by 0.3% on MobileNet-V1. Our results prove that adaptive bit-width is an additional option for adaptive models, which is able to further improve trade-offs between efficiency and accuracy for deep neural networks in addition to previous choices of width, depth, kernel size, and resolution.

### 5.2. Illustration of clipping levels

To understand the impact of S-CL, we visualize the clipping levels from different layers in AB-MobileNet-V1 in

Individual Quantization (SAT)				Adaptive Bit-widths			BitOPs
Name	Bit-width	Size	Top-1 Acc.	Name	Size	Top-1 Acc.	
MobileNet-V1	8 bit	4.10 MB	72.6	AB-MobileNet-V1 [8, 6, 5, 4] bits	4.35 MB	72.4 (-0.2)	36.40 B
MobileNet-V1	6 bit	3.34 MB	72.3			72.4 (0.1)	20.81 B
MobileNet-V1	5 bit	2.96 MB	71.9			72.1 (0.2)	14.68 B
MobileNet-V1	4 bit	2.58 MB	71.3			71.1 (-0.2)	9.67 B
MobileNet-V2	8 bit	3.44 MB	72.5	AB-MobileNet-V2 [8, 6, 5, 4] bits	3.83 MB	72.6 (0.1)	19.25 B
MobileNet-V2	6 bit	2.92 MB	72.3			72.4 (0.1)	11.17 B
MobileNet-V2	5 bit	2.66 MB	72.0			72.1 (0.1)	7.99 B
MobileNet-V2	4 bit	2.40 MB	71.1			70.8 (-0.3)	5.39 B
ResNet-50	4 bit	13.34 MB	76.3	AB-ResNet-50 [4, 3, 2] bits	13.75 MB	76.1 (-0.2)	71.81 B
ResNet-50	3 bit	10.55 MB	75.9			75.8 (-0.1)	43.75 B
ResNet-50	2 bit	7.75 MB	73.3			73.2 (-0.1)	23.71 B

Table 4. Comparison between individual quantization and AdaBits quantization for top-1 validation accuracy (%) of MobileNet-V1/V2 and ResNet-50 on ImageNet.

Fig. 6. We find different bit-widths indeed lead to different values of clipping levels, which generally follow the order that larger bits have relatively larger clipping levels as in the individual models. By privatizing clipping levels to different bit-widths, different optimal values of clipping levels for different bit-widths can be selected and the optimization of the model can be improved.




Figure 6. Clipping levels in different layers from AB-MobileNet-V1.

## 6. Discussion and Future Work

Our approach for adaptive bit-width indicates that bit-width of quantized models is an additional degree of freedom besides channel number, depth, kernel-size and resolution for adaptive models. Previous work [2, 4] demonstrate the possibility to utilize the trained adaptive models for neural architecture search algorithms, based on which improved architectures can be discovered under predefined resource constraints. This suggests we might be able to employ the quantized models with adaptive bit-widths to search for bit-width in each layer or channel, for the purpose of mixed-precision quantization [13, 36, 35, 34, 24]. On the

other hand, adding bit-widths to the list of channel numbers, depth, kernel-size, and resolution for adaptive models gives more freedom to the design of adaptive models, which could enable more powerful adaptive models and facilitate more real world applications.

Evaluation of AdaBits with other quantization methods is another future work. Due to numerous algorithms for neural network quantization, we only select a state-of-the-art algorithm SAT to validate the effectiveness of adaptive bit-width. Since our joint training approach is general and can be combined with any quantization algorithms based on quantization-aware training, we believe similar results can be achieved by combining other quantization approaches with our AdaBits algorithm.

## 7. Conclusion

In this paper, we investigate the possibility to adaptively configure bit-widths for deep neural networks. By comparing models quantized with different bit-widths and experimenting on several baseline methods, we demonstrate the benefits and challenges of quantization with adaptive bit-widths. A joint training approach is proposed for joint optimization of all bit-width in the quantized model. Another treatment named Switchable Clipping Level is also proposed to privatize clipping level parameters to different bit-widths and to eliminate undesired interference between different bit-width. The final AdaBits approach achieves similar predictive accuracies as models quantized with different bit-widths individually, for a wide range of models including MobileNet-V1/V2 and ResNet-50 on the ImageNet dataset. This new kind of adaptive models widen the choices for designing dynamic models which can instantly adapt to different hardware and resource constraints.



## References

- [1] Anonymous. Auto{lr}: A method for automatic tuning of learning rate. In *Submitted to International Conference on Learning Representations*, 2020. under review. 2
- [2] Anonymous. Autoslim: Towards one-shot architecture search for channel numbers. In *Submitted to International Conference on Learning Representations*, 2020. under review. 4, 8
- [3] Anonymous. Rethinking neural network quantization. In *Submitted to International Conference on Learning Representations*, 2020. under review. 1, 2, 3, 4, 5, 7
- [4] Anonymous. Scaling up neural architecture search with big single-stage models. In *Submitted to International Conference on Learning Representations*, 2020. under review. 1, 2, 8
- [5] Yu Bai, Yu-Xiang Wang, and Edo Liberty. Proxquant: Quantized neural networks via proximal operators. *arXiv preprint arXiv:1810.00861*, 2018. 2
- [6] Gabriel Bender, Pieter-Jan Kindermans, Barret Zoph, Vijay Vasudevan, and Quoc Le. Understanding and simplifying one-shot architecture search. In *International Conference on Machine Learning*, pages 549–558, 2018. 2
- [7] Han Cai, Chuang Gan, and Song Han. Once for all: Train one network and specialize it for efficient deployment. *arXiv preprint arXiv:1908.09791*, 2019. 1, 2
- [8] Yukang Chen, Tong Yang, Xiangyu Zhang, Gaofeng Meng, Chunhong Pan, and Jian Sun. Detnas: Backbone search for object detection. *arXiv preprint arXiv:1903.10979*, 2019. 2
- [9] Jungwook Choi, Zhuo Wang, Swagath Venkataramani, Pierce I-Jen Chuang, Vijayalakshmi Srinivasan, and Kailash Gopalakrishnan. Pact: Parameterized clipping activation for quantized neural networks. *arXiv preprint arXiv:1805.06085*, 2018. 2, 3, 6, 7
- [10] Matthieu Courbariaux, Yoshua Bengio, and Jean-Pierre David. Binaryconnect: Training deep neural networks with binary weights during propagations. In *Advances in neural information processing systems*, pages 3123–3131, 2015. 2
- [11] Matthieu Courbariaux, Itay Hubara, Daniel Soudry, Ran El-Yaniv, and Yoshua Bengio. Binarized neural networks: Training deep neural networks with weights and activations constrained to+ 1 or-1. *arXiv preprint arXiv:1602.02830*, 2016. 2
- [12] Ekin D Cubuk, Barret Zoph, Dandelion Mane, Vijay Vasudevan, and Quoc V Le. Autoaugment: Learning augmentation policies from data. *arXiv preprint arXiv:1805.09501*, 2018. 2
- [13] Ahmed T. Elthakeb, Pranjoy Pilligundla, FatemehSadat Mireshghallah, Amir Yazdanbakhsh, Sicun Gao, and Hadi Esmaeilzadeh. Releq: an automatic reinforcement learning approach for deep quantization of neural networks. *arXiv preprint arXiv:1811.01704*, 2018. 2, 8
- [14] Xinyu Gong, Shiyu Chang, Yifan Jiang, and Zhangyang Wang. Autogan: Neural architecture search for generative adversarial networks. *arXiv preprint arXiv:1908.03835*, 2019. 2
- [15] Priya Goyal, Piotr Dollár, Ross Girshick, Pieter Noordhuis, Lukasz Wesolowski, Aapo Kyrola, Andrew Tulloch, Yangqing Jia, and Kaiming He. Accurate, large mini-batch sgd: Training imagenet in 1 hour. *arXiv preprint arXiv:1706.02677*, 2017. 7
- [16] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 770–778, 2016. 1
- [17] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Identity mappings in deep residual networks. In *European conference on computer vision*, pages 630–645. Springer, 2016. 1
- [18] Andrew G Howard, Menglong Zhu, Bo Chen, Dmitry Kalenichenko, Weijun Wang, Tobias Weyand, Marco Andreetto, and Hartwig Adam. Mobilenets: Efficient convolutional neural networks for mobile vision applications. *arXiv preprint arXiv:1704.04861*, 2017. 1, 3
- [19] Benoit Jacob, Skirmantas Kligys, Bo Chen, Menglong Zhu, Matthew Tang, Andrew Howard, Hartwig Adam, and Dmitry Kalenichenko. Quantization and training of neural networks for efficient integer-arithmetic-only inference. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 2704–2713, 2018. 2
- [20] Cong Leng, Zesheng Dou, Hao Li, Shenghuo Zhu, and Rong Jin. Extremely low bit neural network: Squeeze the last bit out with admm. In *Thirty-Second AAAI Conference on Artificial Intelligence*, 2018. 2
- [21] Fengfu Li, Bo Zhang, and Bin Liu. Ternary weight networks. *arXiv preprint arXiv:1605.04711*, 2016. 2
- [22] Chenxi Liu, Liang-Chieh Chen, Florian Schroff, Hartwig Adam, Wei Hua, Alan L Yuille, and Li Fei-Fei. Auto-deeplab: Hierarchical neural architecture search for semantic image segmentation. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 82–92, 2019. 2
- [23] Ilya Loshchilov and Frank Hutter. Sgdr: Stochastic gradient descent with warm restarts. *arXiv preprint arXiv:1608.03983*, 2016. 7
- [24] Qian Lou, Lantao Liu, Minje Kim, and Lei Jiang. Autoqb: Automl for network quantization and binarization on mobile devices. *arXiv preprint arXiv:1902.05690*, 2019. 2, 8
- [25] Naveen Mellempudi, Abhisek Kundu, Dheevatsa Mudigere, Dipankar Das, Bharat Kaul, and Pradeep Dubey. Ternary neural networks with fine-grained quantization. *arXiv preprint arXiv:1705.01462*, 2017. 2
- [26] Asit Mishra and Debbie Marr. Apprentice: Using knowledge distillation techniques to improve low-precision network accuracy. *arXiv preprint arXiv:1711.05852*, 2017. 2
- [27] Asit Mishra, Eriko Nurvitadhi, Jeffrey J Cook, and Debbie Marr. Wrpn: wide reduced-precision networks. *arXiv preprint arXiv:1709.01134*, 2017. 2
- [28] Eunhyeok Park, Junwhan Ahn, and Sungjoo Yoo. Weighted-entropy-based quantization for deep neural networks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 5456–5464, 2017. 2
- [29] Prajit Ramachandran, Barret Zoph, and Quoc V. Le. Searching for activation functions, 2018. 2

- [30] Mohammad Rastegari, Vicente Ordonez, Joseph Redmon, and Ali Farhadi. Xnor-net: Imagenet classification using binary convolutional neural networks. In *European Conference on Computer Vision*, pages 525–542. Springer, 2016. 2
- [31] Mark Sandler, Andrew Howard, Menglong Zhu, Andrey Zhmoginov, and Liang-Chieh Chen. Mobilenetv2: Inverted residuals and linear bottlenecks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 4510–4520, 2018. 1, 3
- [32] Dimitrios Stamoulis, Ruizhou Ding, Di Wang, Dimitrios Lymberopoulos, Bodhi Priyantha, Jie Liu, and Diana Marculescu. Single-path nas: Designing hardware-efficient convnets in less than 4 hours. *arXiv preprint arXiv:1904.02877*, 2019. 2
- [33] Mingxing Tan, Bo Chen, Ruoming Pang, Vijay Vasudevan, Mark Sandler, Andrew Howard, and Quoc V Le. Mnas-net: Platform-aware neural architecture search for mobile. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 2820–2828, 2019. 1
- [34] Stefan Uhlich, Lukas Mauch, Kazuki Yoshiyama, Fabien Cordinaux, Javier Alonso Garcia, Stephen Tiedemann, Thomas Kemp, and Akira Nakamura. Differentiable quantization of deep neural networks. *arXiv preprint arXiv:1905.11452*, 2019. 2, 8
- [35] Kuan Wang, Zhijian Liu, Yujun Lin, Ji Lin, and Song Han. Haq: Hardware-aware automated quantization with mixed precision. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 8612–8620, 2019. 2, 8
- [36] Bichen Wu, Yanghan Wang, Peizhao Zhang, Yuandong Tian, Peter Vajda, and Kurt Keutzer. Mixed precision quantization of convnets via differentiable neural architecture search. *arXiv preprint arXiv:1812.00090*, 2018. 2, 8
- [37] Chen Xu, Jianqiang Yao, Zhouchen Lin, Wenwu Ou, Yuanbin Cao, Zhirong Wang, and Hongbin Zha. Alternating multi-bit quantization for recurrent neural networks. *arXiv preprint arXiv:1802.00150*, 2018. 2
- [38] Jiahui Yu, Linjie Yang, Ning Xu, Jianchao Yang, and Thomas Huang. Slimmable neural networks. *arXiv preprint arXiv:1812.08928*, 2018. 1, 2, 5
- [39] Shuchang Zhou, Yuxin Wu, Zekun Ni, Xinyu Zhou, He Wen, and Yuheng Zou. Dorefa-net: Training low bitwidth convolutional neural networks with low bitwidth gradients. *arXiv preprint arXiv:1606.06160*, 2016. 2, 4
- [40] Shu-Chang Zhou, Yu-Zhi Wang, He Wen, Qin-Yao He, and Yu-Heng Zou. Balanced quantization: An effective and efficient approach to quantized neural networks. *Journal of Computer Science and Technology*, 32(4):667–682, 2017. 2
- [41] Chenzhuo Zhu, Song Han, Huizi Mao, and William J Dally. Trained ternary quantization. *arXiv preprint arXiv:1612.01064*, 2016. 2
- [42] Shilin Zhu, Xin Dong, and Hao Su. Binary ensemble neural network: More bits per network or more networks per bit? In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 4923–4932, 2019. 2
- [43] Barret Zoph and Quoc V Le. Neural architecture search with reinforcement learning. *arXiv preprint arXiv:1611.01578*, 2016. 1, 2
- [44] Barret Zoph, Vijay Vasudevan, Jonathon Shlens, and Quoc V Le. Learning transferable architectures for scalable image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 8697–8710, 2018. 1